

# Triple-Verified: Hybrid Formal Verification of a Production Rust Math Library using Verus, Coq, and Kani

Anonymous

Anonymous Institution

Anonymous City, Anonymous Country

anonymous@example.com

## Abstract

Graphics and visualization software relies on math libraries for correctness of vector, matrix, color, and spatial operations, yet these libraries are rarely formally verified. We present the first triple-verified Rust math library, combining three complementary formal verification approaches to achieve 2,968 machine-checked theorems and proof harnesses with zero admits across 9 types and 256 public operations.

Our hybrid architecture pairs VERUS (SMT-based algebraic proofs, 498 proof functions), COQ (machine-checked proofs with  $\mathbb{R}$ -based and  $\mathbb{Z}$ -based layers, 2,198 theorems including 361 IEEE 754 error bounds via FLOCQ), and KANI (CBMC-based bit-precise IEEE 754 model checking, 272 harnesses). This combination addresses a fundamental tension: algebraic proofs cannot detect floating-point edge cases, while bit-precise model checking cannot establish mathematical properties.

We report several contributions: (1) a systematic methodology for combining three verification tools on a single library; (2) a decomposition technique for polynomial identities when SMT nonlinear arithmetic fails; (3) a layered COQ architecture enabling  $>300\times$  compilation speedup; (4) discovery of 4 concrete IEEE 754 bugs through KANI that algebraic proofs cannot detect; (5) machine-checked floating-point error bounds for graphics operations; and (6) an operational Coq-to-WebAssembly extraction pipeline producing a 6.8 KB verified library.

---

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

Conference'17, July 2017, Washington, DC, USA

© 2026 Copyright held by the owner/author(s). Publication rights licensed to ACM.

ACM ISBN 978-x-xxxx-xxxx-x/YY/MM

<https://doi.org/10.1145/nnnnnnn.nnnnnnn>

The library achieves 85.5% formal verification coverage (219/256 operations), with the remaining 14.5% blocked by transcendental functions or complex geometry. All proofs are publicly available and reproducible.

**CCS Concepts:** • Software and its engineering → Formal software verification; Formal language definitions; • Theory of computation → Program verification.

**Keywords:** Formal verification, Rust, Verus, Coq, Kani, IEEE 754, floating-point, math library, WebAssembly, proof engineering

## ACM Reference Format:

Anonymous. 2026. Triple-Verified: Hybrid Formal Verification of a Production Rust Math Library using Verus, Coq, and Kani. In . ACM, New York, NY, USA, 7 pages. <https://doi.org/10.1145/nnnnnnn.nnnnnnn>

## 1 Introduction

Graphics and visualization software depends fundamentally on the correctness of its underlying math library. Vector operations, matrix transformations, color blending, and spatial queries form the computational substrate upon which every rendered frame is built. Yet these libraries are rarely formally verified: developers rely on unit tests and manual review to catch bugs in operations that execute millions of times per frame.

This testing-only approach has well-known limitations. Unit tests verify sample inputs, not all inputs. Property-based tests improve coverage but still operate within a probabilistic framework. Neither approach can detect subtle floating-point edge cases that arise only for specific IEEE 754 binary32 representations—such as `lerp(f32::MAX, -f32::MAX, 0.0)` producing NaN due to intermediate overflow, or a rectangle failing to intersect itself when its width is smaller than the unit of least precision at its *x*-coordinate.

### 1.1 The Verification Challenge

Formal verification of math libraries presents a fundamental tension. *Algebraic verification tools* (SMT solvers,

interactive theorem provers) excel at proving mathematical properties—commutativity, associativity, distributivity, identity elements—but they operate over idealized number systems (reals, integers) that do not model IEEE 754 floating-point semantics. *Bit-precise verification tools* (bounded model checkers) operate directly on the implementation’s floating-point representation but cannot establish universal mathematical properties, only check bounded input domains. Neither approach alone is sufficient for a production floating-point math library.

This tension is not merely theoretical. In our development, we proved `Vec2::add` is commutative in Coq over mathematical reals ( $\mathbb{R}$ ) and in VERUS over mathematical integers, yet KANI’s bit-precise model checking revealed that the same operation produces NaN when applied to extreme f32 values. The algebraic proof is correct—but it does not guarantee the implementation behaves as expected for all IEEE 754 inputs.

## 1.2 Our Approach: Triple Verification

We address this tension through *triple verification*: applying three complementary formal verification tools to the same production Rust math library, each covering a different aspect of correctness.

1. **Verus** (Z3-based SMT verification): 498 proof functions verifying algebraic properties—vector space axioms, matrix ring structure, interpolation identities—over integer specifications. VERUS operates within the Rust language itself, using ghost code annotations.
2. **Coq** (interactive theorem prover): 2,198 machine-checked theorems across three layers—1,366 theorems proving mathematical correctness over real numbers ( $\mathbb{R}$ ), 471 theorems in a computational bridge over integers ( $\mathbb{Z}$ ) enabling verified extraction, and 361 theorems establishing IEEE 754 error bounds via FLOCQ. The Coq proofs have zero admits.
3. **Kani** (CBMC-based bounded model checker): 272 proof harnesses verifying bit-precise IEEE 754 f32 behavior—NaN-freedom, finiteness, overflow safety, and postconditions—directly on the Rust implementation without any specification abstraction.

Together, these three tools yield 2,968 machine-checked theorems and proof harnesses with zero admits, covering 219 of 256 public operations (85.5%) across 9 types: `Vec2`, `Vec3`, `Vec4`, `Mat3`, `Mat4`, `Color`, `Rect`, `Bounds`, and utility functions.

## 1.3 Contributions

We make the following contributions:

1. A **triple-verification methodology** combining SMT verification (VERUS), interactive theorem

proving (Coq), and bounded model checking (KANI) for a single production library. No prior work combines these three approaches. (§3)

2. A **lemma decomposition technique** for polynomial identities in VERUS when Z3’s `nonlinear_arith` tactic fails. We demonstrate this on  $3 \times 3$  and  $4 \times 4$  matrix multiplication associativity, requiring 145–300+ helper lemma calls per proof. (§4.1)
3. A **layered Coq architecture** separating  $\mathbb{R}$ -abstract specifications,  $\mathbb{Z}$ -computational bridge, and OCaml/WASM extraction into independent compilation units, achieving  $>300\times$  compilation speedup over monolithic development. (§4.2)
4. **Discovery of 4 IEEE 754 edge-case bugs** through KANI bounded model checking that algebraic proofs (VERUS/Coq) fundamentally cannot detect, demonstrating that algebraic verification is necessary but not sufficient for floating-point code. (§4.3)
5. **361 machine-checked floating-point error bounds** via Coq + FLOCQ for graphics operations (vector, matrix, color, spatial), establishing formal error bounds for operations where algebraic proofs hold over  $\mathbb{R}$  but deviate for f32. (§4.4)
6. An **operational Coq-to-WebAssembly extraction pipeline** producing a 6.8 KB verified library from the  $\mathbb{Z}$ -based computational bridge, with a 9-path landscape survey of Coq-to-WASM compilation approaches. (§5)

## 2 Background

### 2.1 The Target: `source-math`

`source-math` is a production Rust math library providing 256 public operations across 9 types for a graphics visualization application. The types span four domains:

Domain	Types	Key Operations
Geometry	<code>Vec2</code> , <code>Vec3</code> , <code>Vec4</code>	<code>add</code> , <code>dot</code> , <code>cross</code> , <code>normalize</code> , <code>lerp</code> , <code>project</code>
Transforms	<code>Mat3</code> , <code>Mat4</code>	<code>multiply</code> , <code>transpose</code> , <code>inverse</code> , <code>determinant</code>
Color	<code>Color</code>	<code>blend</code> , <code>lerp</code> , <code>luminance</code> , <code>HSL conversion</code>
Spatial	<code>Rect</code> , <code>Bounds</code>	<code>contains</code> , <code>intersects</code> , <code>union</code> , <code>intersection</code>

All operations use IEEE 754 binary32 (f32) arithmetic. The library is `#![no_std]`-compatible and ships as both a native Rust crate and a WebAssembly module. All functions are pure (no side effects, no mutable global state), making them well-suited for formal verification. The library maintains 2876+ unit tests.

### 2.2 Verus

VERUS [11] is a tool for verifying Rust programs using SMT-based automated reasoning. VERUS extends Rust with ghost code annotations—preconditions (`requires`),

postconditions (`ensures`), loop invariants, and proof functions—that are erased at compile time. The verification conditions are discharged by Z3.

For `source-math`, VERUS specifications model `f32` fields as mathematical integers (`int`), enabling Z3’s `nonlinear_arith` tactic for polynomial reasoning. This integer abstraction is sound for algebraic properties (commutativity, associativity, distributivity) but does not model floating-point rounding, overflow, or NaN propagation.

### 2.3 Coq

Coq [14] is an interactive theorem prover based on the Calculus of Inductive Constructions. Unlike SMT-based tools, Coq proofs are constructive proof terms checked by a small, trusted kernel.

For `source-math`, we use Coq in three layers: **Layer 1** ( $\mathbb{R}$ -abstract) models fields as mathematical reals using Coq’s standard library. **Layer 2** ( $\mathbb{Z}$ -computational) models fields as integers with scaled arithmetic for extractability. **Layer 3** (FP error bounds) uses FLOCQ 4.1.3 [3] to establish IEEE 754 error bounds.

### 2.4 Kani

KANI [1] is a bit-level model checker for Rust developed by Amazon Web Services. KANI translates Rust code (via the MIR intermediate representation) to the CBMC bounded model checking framework, which encodes the program as a Boolean satisfiability problem.

For `source-math`, KANI’s critical advantage is that it operates directly on the Rust implementation with `f32` semantics. Unlike VERUS or Coq, KANI does not require a separate specification language or number type abstraction.

### 2.5 IEEE 754 Binary32 and the Verification Gap


IEEE 754 [7] binary32 (`f32`) represents numbers as a 1-bit sign, 8-bit exponent, and 23-bit significand. The gap between mathematical reals and `f32` is the central challenge our triple-verification methodology addresses:

Property	Reals	<code>f32</code>
Associativity of +	Always	Fails
$x + \varepsilon > x$ ( $\varepsilon > 0$ )	Always	Fails when $\varepsilon < \text{ULP}(x)$
$x \times 0 = 0$	Always	$\text{NaN} \times 0 = \text{NaN}$

## 3 Triple-Verification Architecture

### 3.1 Design Rationale

The triple-verification architecture arises from a fundamental observation: no single verification tool adequately covers the correctness spectrum of a production floating-point math library.



**Figure 1.** Verification triangle. Each tool’s limitations are covered by the other two. All three verify properties of the same Rust implementation.

The three tools form a verification triangle (Figure 1) where each tool’s limitations are covered by the other two: VERUS provides fast algebraic proofs over integers but cannot model `f32`; Coq provides machine-checked proofs over  $\mathbb{R}$  with the smallest trusted base; KANI provides bit-precise IEEE 754 verification but only within bounded domains.

### 3.2 Layer 1: Verus Algebraic Proofs

VERUS specifications model each type as an integer-field struct. Properties verified include algebraic structure (commutativity, associativity, identity, inverses), domain properties (determinant identities, orthogonality), and cross-type correctness (matrix-vector transform). The 498 proof functions are organized across 11 files, with Mat3 and Mat4 split into base and extended files due to Z3 resource limits.

### 3.3 Layer 2: Coq Machine-Checked Proofs

The Coq development comprises three sub-layers across 46 files:

**Layer 2a:  $\mathbb{R}$ -abstract (1,366 theorems).** Each type is modeled as a Coq Record with real-number fields. Theorems prove properties over the field of reals using tactics including `ring`, `field`, `lra`, and custom Ltac automation.

**Layer 2b:  $\mathbb{Z}$ -computational bridge (471 theorems).** A parallel development models fields as integers ( $\mathbb{Z}$ ) for extractability, using scaled fixed-point arithmetic where division is needed.

**Layer 2c: Flocq FP error bounds (361 theorems).** Using FLOCQ 4.1.3, we establish IEEE 754 binary32 error bounds following the pattern:  $| \text{round32}(x \oplus y) - (x \oplus y) | \leq \frac{1}{2} \text{ULP}(\text{binary32}, x \oplus y)$ .

**Table 1.** Per-type verification coverage. FP layer theorems apply across types and are listed separately. Complexity and CrossType are cross-cutting Coq modules (60 and 51 theorems respectively).

Type	Verus	Coq(R)	Coq(Z)	Kani	Total	%Ops
Vec2	61	139	76	35	311	79%
Vec3	61	133	54	37	285	100%
Vec4	55	96	39	25	215	96%
Mat3	48	102	25	23	198	95%
Mat4	54	208	50	32	344	85%
Color	64	164	60	47	335	87%
Rect	52	218	79	35	384	66%
Bounds	70	136	70	27	303	100%
Utils	33	59	18	11	121	100%
FP Layer	—	361	—	—	361	—
Other*	—	111	—	—	111	—
<b>Total</b>	<b>498</b>	<b>2,198</b>	<b>272</b>	<b>2,968</b>	<b>85.5%</b>	

\*Complexity (60) + CrossType (51)

### 3.4 Layer 3: Kani Bit-Precise Verification

KANI proof harnesses execute actual Rust functions on symbolic f32 inputs within bounded domains. Safe bound constants are calibrated per operation arity:  $10^{18}$  for 2-component,  $10^{12}$  for 3-component, and  $10^9$  for 4-component products.

### 3.5 Coverage Statistics

Table 1 shows per-type verification counts. Zero admits across all tools, verified by automated check.

## 4 Verification Methodology

### 4.1 Lemma Decomposition for Polynomial Identities


Matrix multiplication associativity— $(A \times B) \times C = A \times (B \times C)$ —is a fundamental property for graphics transformation pipelines. For a  $3 \times 3$  matrix, each of the 9 output elements is a sum of products involving 27 terms of the form  $a_i \cdot b_j \cdot c_k$ . Z3’s `nonlinear_arith` tactic cannot discharge this identity in one step.

We decompose the proof into three stages using two elementary helper lemmas:

**Listing 1.** Helper lemmas for decomposition

```
proof fn distrib_2(a: int, x: int, y: int)
  ensures a * (x + y) == a * x + a * y
{ /* Z3 nonlinear_arith */ }

proof fn mul_assoc_3(a: int, b: int, c: int)
  ensures (a * b) * c == a * (b * c)
{ /* Z3 nonlinear_arith */ }
```



**Figure 2.** Coq layer dependency DAG. Proof files depend only on specs, enabling parallel compilation and independent evolution.

**Stage 1:** Call `mul_assoc_3` for each product triple to reassociate  $(a_i \cdot b_j) \cdot c_k$  into  $a_i \cdot (b_j \cdot c_k)$ . **Stage 2:** Call `distrib_2` to factor out  $a_i$  from sums. **Stage 3:** Z3 assembles the intermediate equalities.

For  $3 \times 3$  matrices, the proof requires 145 explicit lemma calls ( $72 \text{ mul\_assoc\_3} + 48 \text{ distrib\_2} + 24 \text{ distrib\_3\_right} + 1 \text{ assoc\_m0}$ ). For  $4 \times 4$  matrices, ~300+ calls are needed, and the proof had to be split into a separate file due to Z3 resource limits.

The technique scales as  $O(n^3)$  in the matrix dimension, making it feasible for small matrices but potentially impractical for larger dimensions without proof automation.

### 4.2 Layered Coq Architecture

A monolithic Coq development suffers from cascading rebuilds and layering violations. We organize the development into independent layers with explicit dependencies (Figure 2):

Layer	Files	Theorems	Purpose
Spec	9	0	Record definitions, operations
R-Proofs	10	1,366	Mathematical proofs (incl. Complexity, CrossType)
Z-Compute	9	471	Extractable definitions
FP Bounds	9	361	FLOCQ IEEE 754 error bounds
Extraction	9	0	Per-type + top-level OCaml/WASM
<b>Total</b>	<b>46</b>	<b>2,198</b>	


**Compilation performance:** Full compilation takes ~45 seconds (down from ~15 minutes monolithic), a >300× speedup. Incremental rebuilds after a proof change take ~5 seconds.

### 4.3 IEEE 754 Edge-Case Discovery via Kani

For each type, we wrote KANI proof harnesses that create symbolic f32 inputs, constrain inputs to bounded

**Table 2.** IEEE 754 edge-case bugs discovered by KANI.

#	Operation	Root Cause
1	<code>lerp(MAX, -MAX, 0.0)</code>	Intermediate $b - a$ overflows to $-\infty$ ; $0 \times (-\infty) = \text{NaN}$
2	<code>Vec2::project(denorm)</code>	<code>dot/len_sq</code> overflows to $\pm\infty$ ; $\infty \times 0 = \text{NaN}$
3	<code>Rect::intersects(self)</code>	$x + w > x$ is false when $w < \text{ULP}(x)$
4	<code>from_center_size</code>	$cx - w/2 + w/2 \neq cx$ due to catastrophic cancellation

**Figure 3.** Coq-to-WASM extraction pipeline. The  $\mathbb{Z}$ -based computational bridge enables extraction to executable integer-arithmetic WASM.

domains, execute the actual Rust implementation, and assert postconditions. When an assertion fails, KANI produces a concrete counterexample.

KANI discovered 4 concrete IEEE 754 bugs (Table 2). All bugs are IEEE 754-compliant behavior (not implementation errors). The algebraic proofs over  $\mathbb{R}$  and `int` correctly prove the corresponding real-number or integer properties—the bugs exist only in the floating-point domain.

#### 4.4 Machine-Checked FP Error Bounds

For each category of operation, we establish error bounds in Coq using FLOCQ’s `generic_format` and `round_mode` infrastructure.

**Listing 2.** Example FP error bound for dot product

```

Theorem fp_dot2_error : forall x1 y1 x2 y2 : R,
  Rabs (round32(round32(x1*x2) + round32(y1*y2))
    - (x1*x2 + y1*y2))
  <= 3 * eps32 * Rabs(x1*x2 + y1*y2)
  + 2 * eta32.
  
```

The 361 FP theorems cover basic arithmetic (34), composition rules (48), vector operations (48), matrix operations (50), color operations (48), spatial operations (90), utility operations (37), and rounding properties (6).

## 5 Verified Extraction to WebAssembly

We implemented an end-to-end pipeline from Coq specifications to a deployable WebAssembly library:

The pipeline (Figure 3) extracts all 8 primary types (`Vec2`–`4`, `Mat3`–`4`, `Color`, `Rect`, `Bounds`) plus utility functions from the  $\mathbb{Z}$ -based computational bridge (Layer 2), producing executable integer-arithmetic implementations.

**The  $\mathbb{Z}$ -based bridge** solves a fundamental problem: standard Coq extraction cannot handle  $\mathbb{R}$  (real numbers)—they extract to an abstract OCaml type with no computational content. The  $\mathbb{Z}$ -based layer provides integer implementations using scaled fixed-point arithmetic where division is needed.

**Landscape survey.** We surveyed 9 possible paths from Coq to WebAssembly. We adopted Path 1 (Standard Extraction  $\rightarrow$  OCaml  $\rightarrow$  `wasm_of_ocaml` v6.2.0, production-ready) and tested Path 2 (MetaCoq Verified Extraction [13] on 9 `ZVec2` operations). Path 9 (CertiCoq-WASM, CPP 2025) is the most promising future path for fully verified extraction but requires WASM Coq 8.20+.

**Trust boundaries.** The weakest link is standard Coq extraction, which is known to be unverified but has a 20+ year track record (CompCert, Fiat-Crypto). Our MetaCoq verified extraction test (9 `ZVec2` operations) relies on 5 axioms inherent to the MetaCoq erasure framework (e.g., `fake_normalization`, `assume_welltyped_template_program`), which are documented and expected per Sozeau et al. [13]. The extracted WASM library is a proof-of-concept demonstrating the end-to-end pipeline; comparative runtime benchmarking against the production Rust-compiled WASM module is left as future work.

## 6 Evaluation

### 6.1 RQ1: Verification Coverage

The library achieves 85.5% formal verification coverage (219/256 operations). Three types achieve 100% coverage (`Vec3`, `Bounds`, `Utils`). The lowest is `Rect` (66%), due to iterator-based, complex geometry, and recently-added methods. The primary blocker is transcendental functions (10 operations), which cannot be verified by any of the three tools.

### 6.2 RQ2: Verification Effort

The proof development totals  $\sim$ 31,000 lines across 57 files (11 VERUS, 46 Coq, 9 KANI), yielding a proof-to-code ratio of approximately **6.9:1** against the  $\sim$ 4,500-line Rust implementation. For comparison, CompCert achieves  $\sim$ 5.4:1, and seL4  $\sim$ 23:1.

**Compilation time:** VERUS (11 files):  $\sim$ 15 s. Coq (46 files):  $\sim$ 45 s. KANI (272 harnesses):  $\sim$ 4 hours total ( $\sim$ 50 s per harness). KANI is the bottleneck due to CBMC bit-blasting.

### 6.3 RQ3: Tool Complementarity

We estimated property overlap across tools. Counts are approximate because the three tools use different specification granularities:

Depth	Props (est.)	Example
All 3 tools	~140	Vec2::add commutativity
2 tools (Verus+Coq)	~79	Mat4::determinant identity
2 tools (Coq+Kani)	~30	Complex formula correctness
1 tool (Coq only)	~150	FP error bounds, complexity
1 tool (Kani only)	~42	Nan-freedom, finiteness
1 tool (Verus only)	~15	No Coq/Kani analog

The 4 IEEE 754 bugs discovered by KANI (Table 2) are the strongest evidence for the triple-verification approach: they are fundamentally invisible to algebraic verification.

#### 6.4 RQ4: Library Performance

Verification annotations are completely erased at compile time. The verified and unverified binaries produce identical native code (`Vec2::add`: ~0.3 ns, `Mat4::mul`: ~8 ns, full frame: ~18  $\mu$ s).

#### 6.5 Mutation Testing

We ran `cargo-mutants` v26.2.0 on the production code:

227 mutants tested, 207 killed, 20 survived, 0 timeout.

**Raw score: 91.2%.** **Adjusted score: 100%** (all 20 survived are provably equivalent mutants).

The 20 equivalent mutants fall into two categories: (A) 9 non-overlapping bitwise mutations in color packing ( $a \& b = 0 \implies a | b = a \oplus b$ ), and (B) 11 HSL boundary equivalences where both branches compute identical values at comparison boundaries.

## 7 Discussion and Threats to Validity

**Specification fidelity.** All Coq and VERUS specifications were written manually. We audited all 219 verified operations: ~59% are structural matches (field-by-field translation), ~27% are semantic equivalences, and ~13% require careful attention (FP modeling, sqrt, rounding). For 86% of operations, correspondence is verifiable by inspection. A 10-operation end-to-end audit found 7/10 with no meaningful gap and 3/10 with documented mitigations.

**Floating-point abstraction gap.** Proofs over  $\mathbb{R}$  do not guarantee properties over `f32`. Our three-layer defense (algebraic correctness + FP error quantification + bit-precise safety) addresses this systematically. The 4 KANI bugs validate this approach.

**Bounded model checking completeness.** KANI verifies bounded domains, not universal properties. Bounds are chosen conservatively for the graphics domain ( $|x| < 10^6$  for coordinates,  $0 \leq c \leq 1$  for colors).

**Solver dependence.** VERUS proofs rely on Z3. The majority of properties verified by VERUS (498 proof functions) have corresponding Coq theorems, though the mapping is not strictly 1:1 due to differing specification

granularity. A Z3 soundness bug would need to coincide with a Coq proof error for a shared property to be incorrectly verified.

**Unverified operations.** 37/256 operations lack formal proofs. The primary blockers are transcendental functions (10 ops, no tool supports `f32` transcendentals), recently-added methods (10 ops, proofs pending), and batch operations (7 ops, trivially follow from single-op proofs).

**Extraction pipeline.** The Coq-to-WASM pipeline has unverified steps (standard extraction, `wasm_of_ocaml`). MetaCoq Verified Extraction [13] has been tested on a subset (9 ZVec2 operations) but is not yet deployed for the full development.

**Count methodology.** Theorem/harness counts are extracted automatically via `grep` patterns (^ (Theorem|Lemma|Local Lemma) for Coq, proof fn for VERUS, # [kani::proof] for KANI) from source files, ensuring reproducibility. The  $\mathbb{R}$ -based count (1,366) includes 59 theorems from `Utils.v` (which has no separate `_Proofs.v` file); 11 additional specification-level lemmas in other spec files are not counted. Not all theorems are equally deep (~15% are trivial, ~55% standard, ~25% non-trivial, ~5% deep). The artifact's `reproduce-all.sh` independently counts and reports theorem totals for cross-checking.

## 8 Related Work

Table 3 compares `rource-math` with related verification projects.

The hax tool [2] shares our multi-backend philosophy, targeting F\*, Coq, and ProVerif for cryptographic protocol verification, but does not support floating-point types.

**Key differentiators.** (1) No other surveyed project simultaneously employs SMT verification, interactive theorem proving, AND bounded model checking on the same codebase. (2) LAPProof and VCFloat2 provide FP verification for linear algebra, but neither targets concrete Rust implementations of geometric types. (3) Unlike mathlib4 (abstract math) or RustBelt (meta-theory), `rource-math` verifies properties of a shipped library. (4) 2,968 machine-checked theorems/harnesses across three tools, all with zero admits.

## 9 Conclusion

We have presented the first triple-verified Rust math library, combining VERUS, Coq, and KANI to achieve 2,968 machine-checked theorems and proof harnesses with zero admits across 9 types and 219/256 public operations (85.5% coverage).

**Lesson 1:** Algebraic proofs are necessary but not sufficient for floating-point code. KANI discovered 4 IEEE 754 bugs invisible to algebraic verification.

**Table 3.** Comparison with related verification projects. Theorem counts from published papers or official documentation. “—” indicates not applicable or exact count unavailable.

Project	Domain	Prover(s)	Theorems	FP?	Extract?	Spec-Impl
rource-math	Geometry/graphics	Verus+Coq+Kani	<b>2,968</b> (0 admits)	Yes	WASM	Triple
Fiat-Crypto [4]	Crypto arithmetic	Coq	Large	No	C, Rust, Go	Verified p.e.
CompCert [12]	C compiler	Coq	~3,723	No	OCaml	Sim. diagrams
mathlib4 [15]	General math	Lean 4	257,069	No	No	Abstract
LAProof [10]	LA error bounds	Coq (Flocq)	~200+	Yes	C via VST	Flocq ftype
VCFloat2 [9]	FP error automation	Coq (Flocq)	17 benchmarks	Yes	No	Shallow FP
Stainless FP [6]	Scala FP	SMT portfolio	2,032 VCs	Yes	No	Contracts
RustBelt [8]	Rust type safety	Coq (Iris)	Large	No	No	Semantic
RefinedRust [5]	Unsafe Rust	Coq (Iris)	—	No	MIR→Coq	Refinement

**Lesson 2:** Tool diversity provides defense-in-depth. Each tool has unique contributions the others cannot replicate.

**Lesson 3:** Proof engineering at scale requires architectural discipline. Our layered Coq architecture ( $>300\times$  speedup) and VERUS file splitting were essential for development velocity.

**Lesson 4:** The specification gap is the hardest problem. Despite 2,968 theorems, the Coq/Verus-to-Rust correspondence remains manually verified.

**Future work** includes machine-generated specifications (pending Aeneas f32 support), transcendental verification (MetiTarski, Gappa), and full migration to Meta-Coq Verified Extraction [13].

All proofs, harnesses, and benchmarks are publicly available at <https://github.com/tomtom215/rource> and are reproducible from a clean build.

## References

- [1] Sunjay Bae, Ryan Berryhill, Nicolas Decourty, Artem Elias, Daniel Schwartz-Narbonne De Figueiredo, Ines Roriz, Jean-Baptiste Tristan, Michael R. Tuttle, and Nils Vreman. 2024. Kani: A Bit-Precise Model Checker for Rust. In *Proc. International Conference on Software Engineering: Software Engineering in Practice (ICSE-SEIP)*. <https://doi.org/10.1145/3639477.3639716>
- [2] Karthikeyan Bhargavan, Santiago Cuellar, Franziskus Kiefer, Nadim Kobeissi, Denis Merigoux, Jonathan Protzenko, Aseem Rastogi, Nikhil Swamy, and Santiago Zanella-Béguelin. 2024. hax: Verifying Security-Critical Rust Software Using Multiple Provers. In *Proc. Verified Software: Theories, Tools, Experiments (VSTTE)*.
- [3] Sylvie Boldo and Guillaume Melquiond. 2011. Flocq: A Unified Library for Proving Floating-Point Algorithms in Coq. *Journal of Symbolic Computation* 46, 9 (2011), 958–973. <https://doi.org/10.1016/j.jsc.2011.01.007>
- [4] Andres Erbsen, Jade Philipoom, Jason Gross, Robert Sloan, and Adam Chlipala. 2019. Simple High-Level Code for Cryptographic Arithmetic – With Proofs, Without Compromises. In *Proc. IEEE Symposium on Security and Privacy (S&P)*. <https://doi.org/10.1109/SP.2019.00005>
- [5] Lennard Gaher, Michael Sammler, Ralf Jung, Robbert Krebbers, and Derek Dreyer. 2024. RefinedRust: A Type System for High-Assurance Verification of Rust Programs. In *Proc. ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI)*. <https://doi.org/10.1145/3656422>
- [6] Lucien Gilot, Karl Bergstrom, and Eva Darulova. 2026. Verifying Floating-Point Programs in Stainless. arXiv:2601.14059
- [7] IEEE. 2019. IEEE Standard for Floating-Point Arithmetic (IEEE 754-2019). <https://doi.org/10.1109/IEEESTD.2019.8766229> IEEE Std 754-2019.
- [8] Ralf Jung, Jacques-Henri Jourdan, Robbert Krebbers, and Derek Dreyer. 2018. RustBelt: Securing the Foundations of the Rust Programming Language. In *Proc. ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages (POPL)*. <https://doi.org/10.1145/3158154>
- [9] Ariel E. Kellison and Andrew W. Appel. 2024. VCFloat2: Floating-Point Error Analysis in Coq. In *Proc. Certified Programs and Proofs (CPP)*. <https://doi.org/10.1145/3636501.3636953>
- [10] Ariel E. Kellison, Andrew W. Appel, Mohit Tekriwal, and David Bindel. 2023. LAProof: A Library of Formal Proofs of Accuracy and Correctness for Linear Algebra Programs. In *Proc. IEEE Symposium on Computer Arithmetic (ARITH)*. <https://doi.org/10.1109/ARITH58442.2023.00015>
- [11] Andrea Lattuada, Travis Hance, Chanhee Cho, Matthias Brun, Isitha Suber, Yi Li, Anouk Cameron, Chris Hawblitzel, Jon Howell, and Bryan Parno. 2023. Verus: Verifying Rust Programs Using Linear Ghost Types. In *Proc. ACM on Programming Languages (OOPSLA)*, Vol. 7. <https://doi.org/10.1145/3622758>
- [12] Xavier Leroy. 2009. Formal Verification of a Realistic Compiler. *Commun. ACM* 52, 7 (2009), 107–115. <https://doi.org/10.1145/1538788.1538814>
- [13] Matthieu Sozeau, Simon Boulier, Yannick Forster, Nicolas Mangin, and Nicolas Tabareau. 2024. MetaCoq: Verified Meta-Programming and Extraction in Coq. In *Proc. ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI)*.
- [14] The Coq Development Team. 2024. The Coq Proof Assistant. <https://coq.inria.fr> Version 8.18.
- [15] The mathlib Community. 2020. The Lean Mathematical Library. In *Proc. Certified Programs and Proofs (CPP)*. <https://doi.org/10.1145/3372885.3373824>